

KUBERNETES

Kubernetes (often abbreviated as **K8s**) is an open-source platform for **automating the deployment, scaling, and management of containerized applications**. It helps run applications reliably across clusters of machines.

Structure of Kubernetes consists of a **master-node (control plane)** and multiple **worker nodes**.

1. Control Plane (Master Node)

Responsible for managing the entire Kubernetes cluster.

- **API Server:** Entry point for all commands (kubectl or API calls).
- **Scheduler:** Assigns workloads (pods) to nodes based on resources and rules.
- **Controller Manager:** Ensures cluster state matches the desired configuration (e.g., restarts failed pods).
- **etcd:** A distributed key-value store to save cluster configuration and state.
- **Cloud Controller Manager (optional):** Integrates with cloud provider APIs.

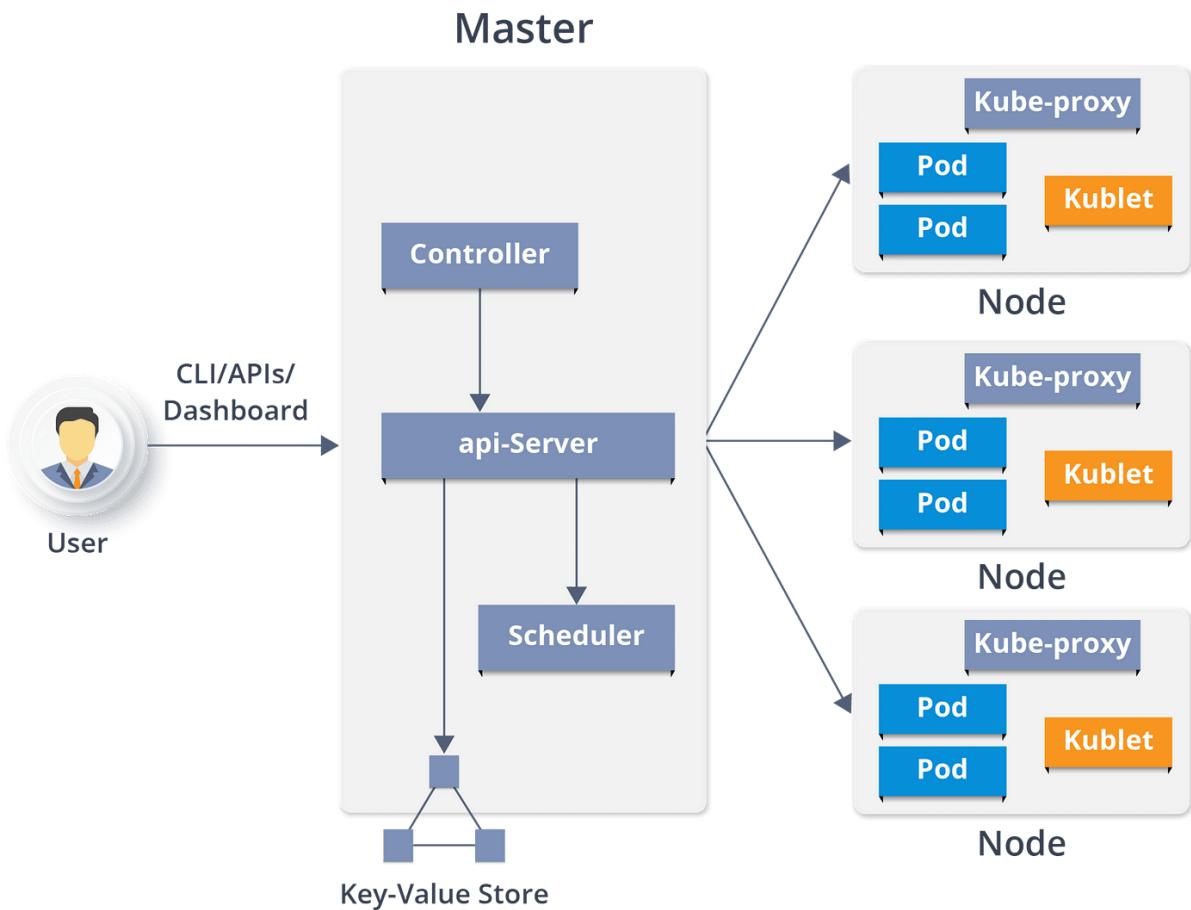
2. Worker Nodes

Run the actual applications (containers).

- **Kubelet:** Communicates with the control plane, manages containers on the node.
 - **Kube Proxy:** Manages networking and load balancing for pods.
 - **Container Runtime:** Runs containers (e.g., Docker, containerd).
-

3. Key Concepts

- **Pods:** The smallest deployable unit; contains one or more containers.
- **Services:** Expose pods to the network; handles load balancing.
- **Deployments:** Manage the rollout and scaling of pods.
- **Namespaces:** Virtual clusters within a Kubernetes cluster for isolation.



PODS

A **Pod** is the **smallest and simplest deployable unit** in Kubernetes. It represents a **single instance of a running process** in your cluster.

Key Features of a Pod:

1. One or More Containers

Most Pods run a **single container**, but they can run multiple containers that need to share:

- **Storage volumes**
- **Network (same IP address & port space)**

2. Shared Resources

- **Network namespace**: All containers in a Pod share the same IP and port range.
- **Storage volumes**: Shared persistent or ephemeral storage.

3. Lifecycle

- Pods are **ephemeral**: If a Pod dies, it's replaced by a new one (with a different ID).
- Higher-level objects like **Deployments** or **StatefulSets** manage this replacement.

Example Use Cases:

- Running a **web server** container.
 - A Pod with a main app container and a **sidecar** container (e.g., for logging or proxying).
-

Analogy:

A Pod is like a **wrapper** around one or more containers that are tightly coupled and should always be scheduled on the same machine.

Deployment

Deployment – [Deployment + Replica Set + Pod]

- **Purpose**: Manages stateless applications.
- **Use Case**: Automatically handles rolling updates, rollbacks, and scaling.

- **Behavior:** Creates **ReplicaSets** underneath to ensure a specified number of pod replicas.

Example: Running multiple instances of a web server like Nginx.

ReplicaSet

ReplicaSet [No. of Pods = No of Replica Set]

- **Purpose:** Ensures a specified number of pod replicas are always running.
- **Use Case:** Low-level component, typically not used directly—used by **Deployments**.
- **Behavior:** Watches pods and replaces them if they fail or are deleted.

Example: A Deployment creates a ReplicaSet which keeps 3 copies of a pod alive.

StatefulSet

StatefulSet

- **Purpose:** Manages **stateful applications** that need **persistent storage** and **stable network identity**.
- **Use Case:** Databases like MySQL, MongoDB, Cassandra.
- **Behavior:**
 - Pods get **unique, stable names** (e.g., mysql-0, mysql-1)
 - Supports ordered deployment, scaling, and updates.
 - Each pod can retain its **persistent volume** even after being deleted.

DaemonSet

- No of Pods = No of Worker Nodes

A **DaemonSet** ensures that a **specific Pod runs on all (or selected) nodes** in your Kubernetes cluster.

Key Features:

1. Runs One Pod per Node

- Automatically adds the Pod to **every new node** added to the cluster.
- Removes the Pod when a node is removed.

2. Use Cases:

- **Log collection agents** (e.g., Fluentd, Filebeat)
- **Monitoring agents** (e.g., Prometheus Node Exporter)
- **Security tools or network plugins**

3. No Replica Count Needed

Unlike Deployments, you don't specify replicas. It runs exactly one Pod **per node** (unless restricted by selectors or taints/tolerations).

Analogy:

Think of a DaemonSet like a **background service** that runs on every machine in your cluster.

Step-by-Step Process:

1. You define the desired state

You create a YAML or JSON file describing what you want (e.g., "run 3 replicas of a web server") and submit it using kubectl or an API call.

2. API Server receives the request

The **API Server** accepts and validates your request, then records the desired state in **etcd** (the cluster's database).

3. Scheduler assigns Pods to Nodes

The **Scheduler** looks for a suitable node with enough resources and assigns the Pod to it.

4. Kubelet creates the Pod on the Node

The **Kubelet** on that node reads the desired state, pulls the container

image, and starts the container using the **container runtime** (like Docker or containerd).

5. Kube Proxy sets up networking

The **Kube Proxy** ensures the Pod can communicate with other Pods and Services inside or outside the cluster.

6. Controller Manager ensures desired state

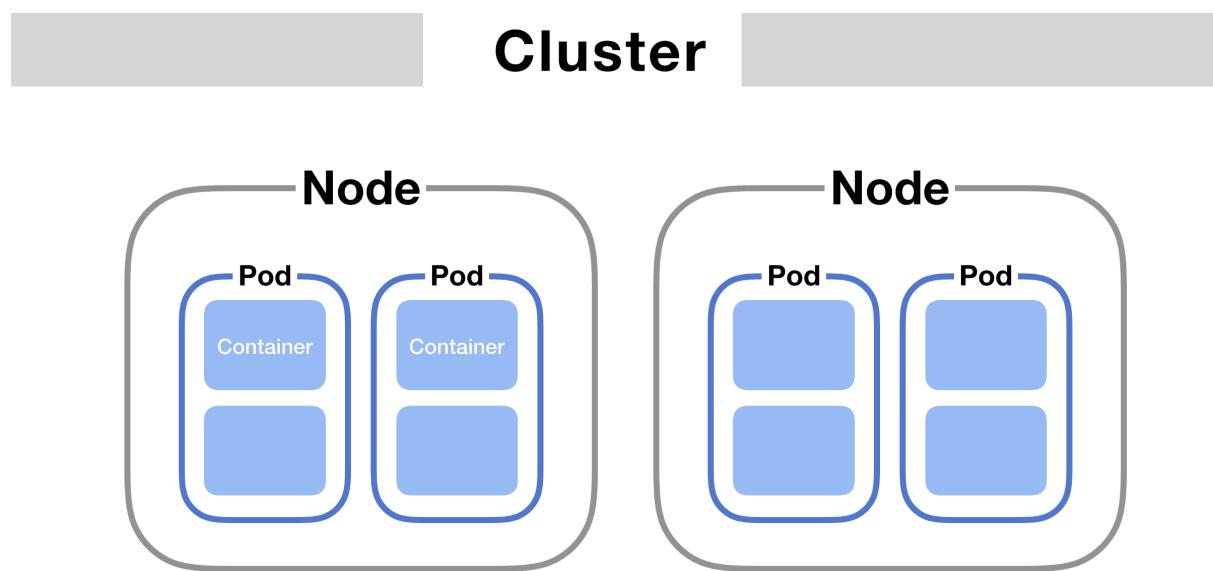
If a Pod crashes or a node fails, the **Controller Manager** detects it and recreates the Pod to match the desired state.

Kubernetes = Continuous Reconciliation

Kubernetes constantly monitors the current state and **automatically makes changes** to match the desired state you defined.

Diagrams and Sample Yaml Files

PODS



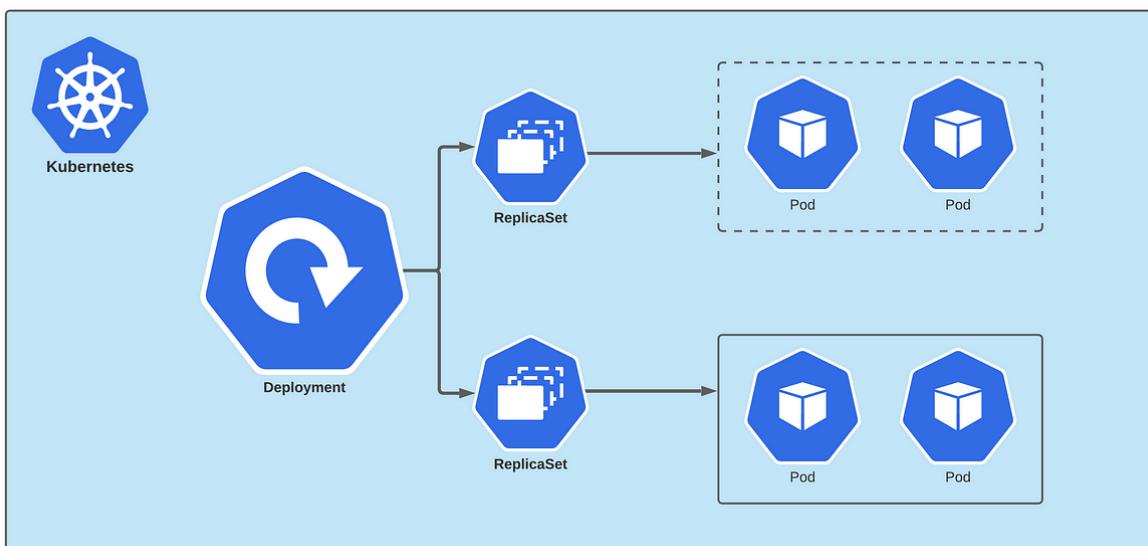
YAML FILE:

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: nginx
spec:
  containers:
  - name: nginx-container
    image: nginx:latest
    ports:
    - containerPort: 80
```

Explanation:

- **apiVersion:** Kubernetes API version to use (v1 for core resources like Pods).
- **kind:** Type of object (Pod).
- **metadata:** Name and labels for the Pod.
- **spec:** Specification of the containers inside the Pod.
- **containers:**
 - **name:** Name of the container.
 - **image:** Docker image to run (nginx:latest).
 - **ports:** Container port to expose (80 for web traffic).

Deployment



YAML FILE

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
  labels:
    app: nginx
spec:
  replicas: 3 # Number of pod replicas
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
```

Explanation:

- **replicas: 3** — Runs 3 identical Pods for high availability.
- **selector.matchLabels** — Matches Pods with label app: nginx to manage them.
- **template** — Defines the Pod template (just like a regular Pod YAML).
- **image: nginx:latest** — Pulls the latest official Nginx image.

ReplicaSet

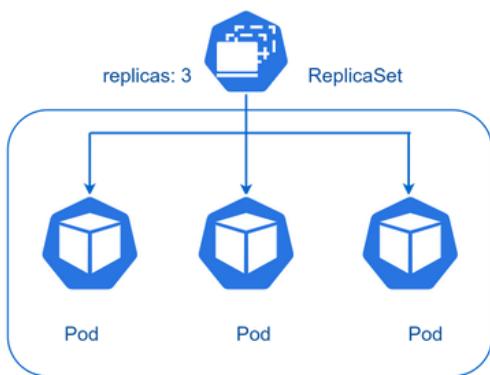


Fig: ReplicaSet

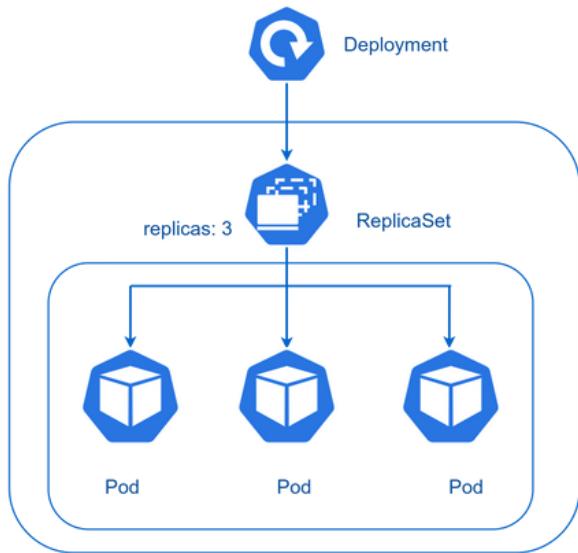
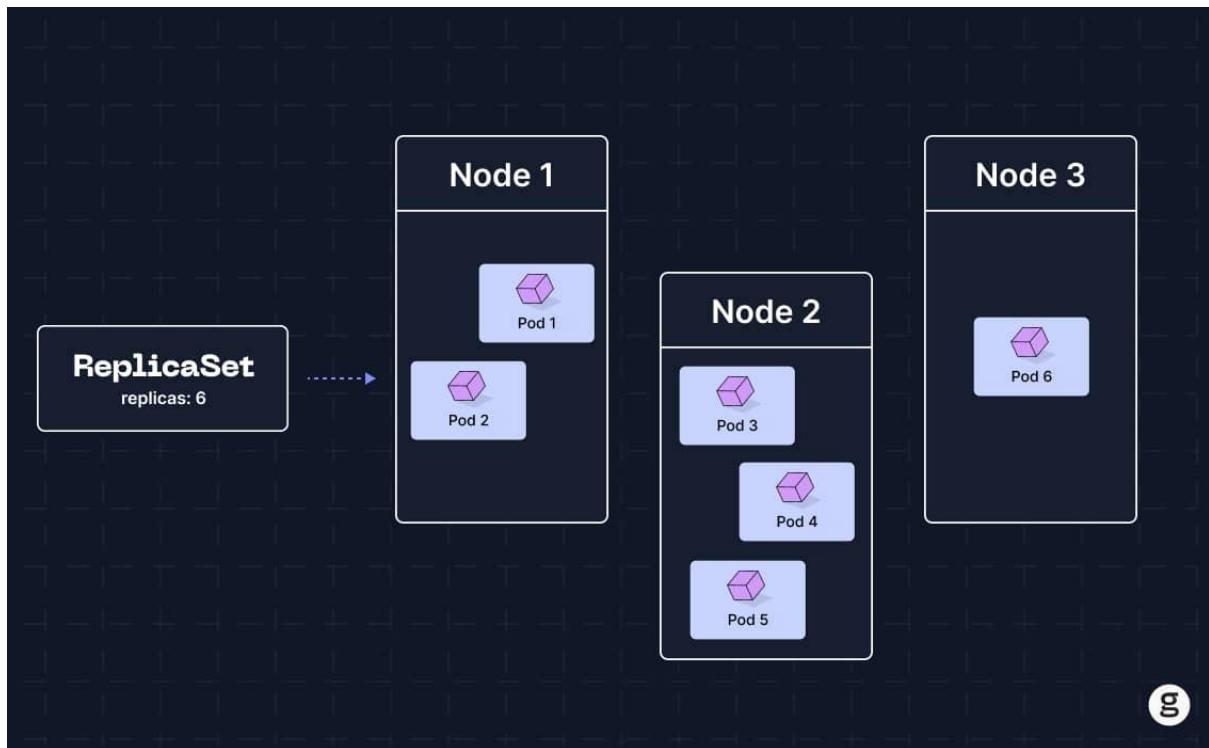


Fig: Deployment

#goglides



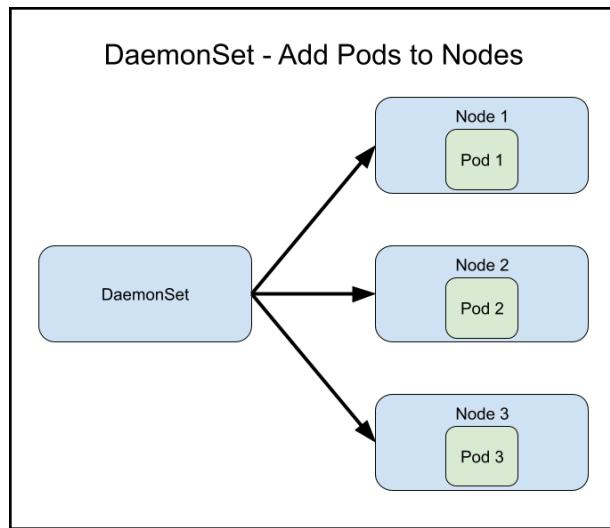
YAML FILE

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-replicaset
  labels:
    app: nginx
spec:
  replicas: 3 # Desired number of Pods
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
          ports:
            - containerPort: 80
```

Explanation:

- **kind: ReplicaSet** — Tells Kubernetes to manage Pods at the replica level.
- **replicas: 3** — Runs 3 instances of the Pod.
- **selector.matchLabels** — Must match the labels in the Pod template exactly.
- **template** — Pod specification just like a standalone Pod YAML

DaemonSet

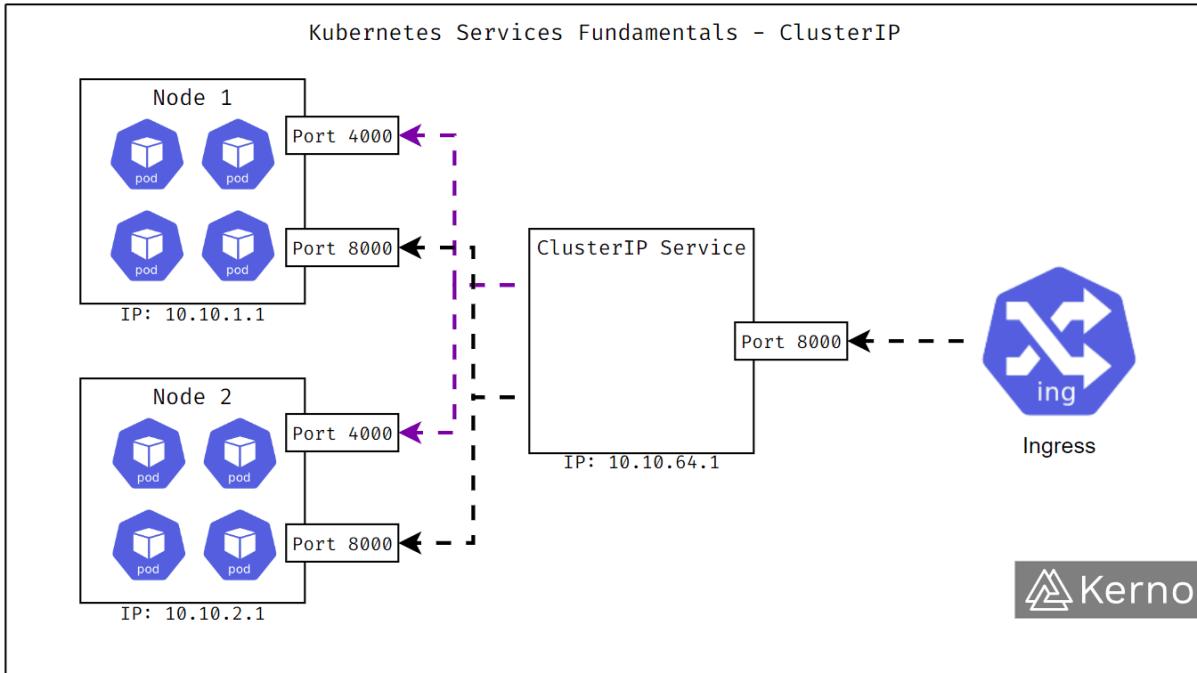


YAML FILE

```
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: nginx-daemonset
  labels:
    app: nginx
spec:
  selector:
    matchLabels:
      name: nginx
  template:
    metadata:
      labels:
        name: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:latest
      ports:
        - containerPort: 80
```

Explanation:

- **kind: DaemonSet** — Ensures a Pod is scheduled on every node.
- **selector.matchLabels** — Matches pods labeled with name: nginx.
- **template** — Pod template to run on each node.
- **image: nginx:latest** — Runs the official Nginx container.



ClusterIP

1. Exposes the Service **internally** within the Kubernetes cluster.
2. Assigns a **virtual IP (VIP)** that can only be accessed from inside the cluster.
3. Internal communication between microservices (e.g., frontend → backend).
4. A backend API accessed only by other services inside the cluster.

NodePort

1. Exposes the Service on a **static port (30000–32767)** on each Node's IP.
2. Maps a port on every node to your Service; you access it via NodeIP:NodePort.
3. Development, simple external access without a load balancer.
4. Access your app at <http://<node-ip>:30080>.

LoadBalancer

1. Creates an **external load balancer** using a cloud provider (AWS, Azure, GCP).
2. Assigns a **public IP** that routes traffic to your Service.
3. Production workloads needing public access.
4. Serving your app to end users via the internet.

ExternalName

1. Maps a Service name to an **external DNS name**.
2. Returns a CNAME record instead of routing traffic through the cluster.
3. Accessing external services (e.g., SaaS APIs) from within Kubernetes.
4. Internally using my-sql-service.default.svc.cluster.local to reach db.example.com.

YAML FILES

ClusterIP

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-clusterip
labels:
  app: nginx
spec:
  selector:
    app: nginx # This must match the
               # pod/deployment label
  ports:
    - protocol: TCP
      port: 80      # Port to expose the
                   # service
      targetPort: 80 # Port on the Pod
                   # container
    type: ClusterIP # Default type (can
                     # omit this line)
```

☞ **SELECTOR.APP:** NGINX MATCHES THE PODS FROM YOUR DEPLOYMENT (MUST MATCH LABELS).

☞ **PORT** IS HOW OTHER SERVICES INSIDE THE CLUSTER WILL CONNECT TO IT.

☞ **TARGETPORT** IS THE PORT YOUR NGINX CONTAINER IS LISTENING ON.

NodePort

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-nodeport
  labels:
    app: nginx
spec:
  type: NodePort      # Expose service via a port on each node
  selector:
    app: nginx        # Must match labels of the target pods
  ports:
    - protocol: TCP
      port: 80          # Internal service port
      targetPort: 80     # Container port in the pod
      nodePort: 30080    # Optional: static port (range 30000–32767)
```

- **ACCESS THE APP AT:**
HTTP://<NODE-IP>:30080
- **PORT:** CLIENTS INSIDE THE CLUSTER USE THIS.
- **NODEPORT:** EXTERNAL CLIENTS USE THIS TO REACH THE SERVICE.
- **TARGETPORT:** THE ACTUAL CONTAINER PORT INSIDE THE POD.

IF YOU OMIT NODEPORT, KUBERNETES WILL AUTO-ASSIGN ONE FROM THE 30000–32767 RANGE.

LoadBalancer

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-loadbalancer
labels:
  app: nginx
spec:
  type: LoadBalancer      # Requests
  an external load balancer
  selector:
    app: nginx           # Must match
    the Deployment/Pod labels
  ports:
    - protocol: TCP
      port: 80            # Exposed service
      port
      targetPort: 80       # Pod/container
      port
```

- ❑ CREATES A PUBLIC IP AND LOAD BALANCER (ON AWS, GCP, AZURE, ETC.).
- ❑ TRAFFIC TO THE EXTERNAL IP IS LOAD-BALANCED TO THE BACKEND PODS.
- ❑ THE SERVICE TARGETS PODS WITH APP: NGINX LABELS.