

# Kubernetes From Scratch – Complete Guide

Welcome to this comprehensive, visually enriched guide to Kubernetes (K8s)! This guide uses help you understand container orchestration step-by-step.

Let's orchestrate some magic! 

---

## What is Kubernetes?

Kubernetes is an **open-source container orchestration platform** that automates the deployment, scaling, and management of containerized applications across clusters. Originally developed by Google, it is now maintained by the **Cloud Native Computing Foundation (CNCF)**.

## Key Benefits

<u>Feature</u>	<u>Description</u>
<b>High Availability</b>	Ensures applications run redundantly across nodes to reduce downtime.
<b>Zero-Downtime Deployments</b>	Performs rolling updates without service interruption.
<b>Auto-Scaling</b>	Adjusts resources automatically based on metrics (CPU, memory).
<b>Auto-Healing</b>	Restarts failed containers, reschedules pods, and maintains replicas.

## What Does Container Orchestration Do?

Kubernetes handles the heavy lifting of managing containers at scale.

### Core Functions

- **Provision & Deploy** containers across multiple nodes
- **High Availability** through balanced distribution
- **Auto-Scaling** workloads (via HPA, VPA)
- **Efficient Resource Allocation** (CPU, memory, storage)
- **Load Balancing & Traffic Routing**
- **Service Discovery** between microservices
- **Secure Communication** via Network Policies

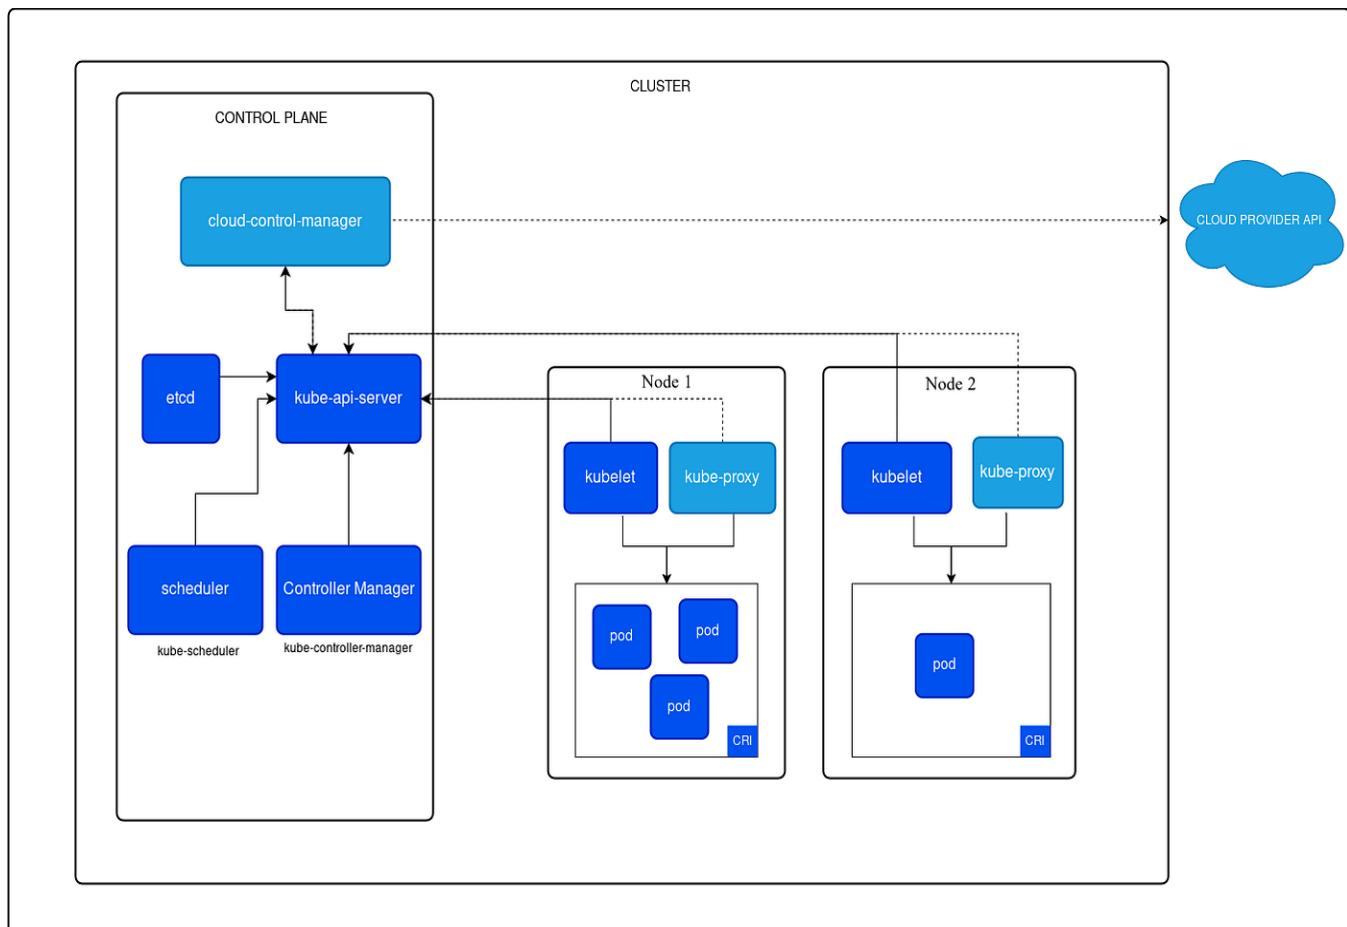
# Popular Orchestrators Comparison

Orchestrator	Strengths	Best For	Drawbacks
Kubernetes	Feature-rich, flexible, huge ecosystem	Enterprise apps	Steep learning curve
Docker Swarm	Simple, Docker-native	Small teams, quick deployments	Less advanced
Apache Mesos	Handles mixed workloads	Big data clusters	Complex, declining adoption

## Kubernetes Components & Architecture

Kubernetes operates through a cluster — a group of machines working together.

## Cluster Overview



**Master Node = Control plane = Brain**

**Worker Nodes = Muscle**

**Multi-Master Setup = High Availability**

## Master Node Components

Component	Role	Key Features
API Server	Central endpoint for all REST/K8s requests	Handles kubectl, validates config, exposes /metrics
Controller Manager	Maintains desired state	Includes node/replica controllers
Scheduler	Assigns pods to nodes	Considers resources, affinities, taints
etcd	Distributed key-value store	Stores entire cluster state

### Essential Add-ons

- **CoreDNS** – Service discovery
  - **Kubernetes Dashboard** – Web UI
  - **Container Runtime** – Docker, containerd, CRI-O
- 

## Worker Node Components

<u>Component</u>	<u>Role</u>	<u>Key Features</u>
Kubelet	Node agent managing pod lifecycle	Runs containers, performs health checks
Kube-Proxy	Network routing for services	Implements load balancing via iptables/IPVS

*Pro Tip: Kubelet communicates only with the API Server, never directly with other nodes.*

---

## kubectl — The Swiss Army Knife CLI

General syntax:

```
kubectl [command] [type] [name] [flags]
```

### Examples of Commands

- **Commands:** get, apply, create, delete, describe
- **Types:** pod, deployment, service, namespace
- **Flags:** -n, -o yaml, --all-namespaces

## kubectl Cheat Sheet

Action	Command Example
List resources	kubectl get pods -A
Describe	kubectl describe pod my-pod
Create/Apply	kubectl apply -f deployment.yaml
Scale	kubectl scale deploy/my-app --replicas=5
Logs & Exec	kubectl logs my-pod kubectl exec -it my-pod -- /bin/sh
Delete	Kubectl delete pod my-pod
Port-forward	kubectl port-forward svc/my-svc 8080:80
Events	kubectl get events --sort-by=.metadata.creationTimestamp

---

## Core Kubernetes Concepts

### Pod — The Smallest Unit

Pods wrap around one or more containers that share:

- Network namespace (one IP)
- Storage volumes
- Lifecycle management rules

### Pod Types

- Single-container Pod
  - Multi-container Pod (Sidecar pattern)
- 

## Deployment — Declarative Scaling

- Manages ReplicaSets
- Handles rolling updates/rollbacks
- Ensures stateless workloads scale

### Deployment YAML

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-app
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-app
  template:
    metadata:
      labels:
        app: my-app
    spec:
      containers:
        - name: nginx
          image: nginx:1.14
```

## Service — Stable Networking

<u>Type</u>	<u>Use Case</u>	<u>Exposure</u>
ClusterIP	Internal-only	Virtual IP inside cluster
NodePort	External access via nodes	<NodeIP>:<NodePort>
LoadBalancer	Cloud LB	Public external IP

---

## Namespace — Virtual Clusters

Used for isolation, organization, and resource limits.

Default namespaces:

- default
  - kube-system
  - kube-public
- 

## Secrets, ConfigMaps & Storage

<u>Concept</u>	<u>Purpose</u>	<u>Security Notes</u>
Secrets	Store sensitive info	Base64 encoded, not encrypted
ConfigMaps	Store config data	Mounted as env or volumes
Volumes	Persist data	Supports emptyDir, hostPath, PVC

---

## ReplicaSet & DaemonSet

- **ReplicaSet** – Maintains N identical Pods
  - **DaemonSet** – Runs one Pod per node (logs, monitoring)
- 

## Labels, Selectors & Annotations

- **Labels:** metadata to organize resources
  - **Selectors:** query/filter resources
  - **Annotations:** metadata for external tools
- 

## Affinity, Anti-Affinity & Taints

Control where Pods are scheduled:

- **Node Affinity** – Match node labels
- **Pod Affinity/Anti-Affinity** – Group or separate pods
- **Taints/Tolerations** – Restrict scheduling on specific nodes

## RBAC — Role-Based Access Control

<u>Resource</u>	<u>Scope</u>	<u>Purpose</u>
Role	Namespace	Defines permissions
ClusterRole	Cluster	Global permissions
RoleBinding	Namespace	Binds Role
ClusterRoleBinding	Cluster	Binds ClusterRole

**Best Practice:** Always follow **least privilege**.

---

## Networking & Ingress

### Ingress — Smart HTTP Routing

Feature	Benefit
Path-based routing /api → backend1, /web → frontend	
TLS termination	HTTPS security
Host-based routing api.example.com, app.example.com	

---

## Scaling & Autoscaling

### Horizontal Pod Autoscaler (HPA)

Automatically scales deployments.

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: my-app-hpa
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: my-app
  minReplicas: 1
  maxReplicas: 10
  metrics:
  - type: Resource
    resource:
      name: cpu
      target:
        type: Utilization
        averageUtilization: 50
```

### Vertical Pod Autoscaler (VPA)

Adjusts CPU/memory requests automatically (may restart pods).

# Monitoring, Logging & Troubleshooting

## Built-in Tools

- `kubectl get events -w`
- `kubectl logs my-pod -f`
- `kubectl top pods`

## Observability Stack

<u>Tool</u>	<u>Purpose</u>
<b>Prometheus</b>	Metrics
<b>Grafana</b>	Dashboards
<b>EFK Stack</b>	Centralized logs
<b>Jaeger</b>	Tracing

Troubleshooting Flow:

- `kubectl describe pod <name>`
  - `kubectl exec -it <pod> -- curl <service>`
- 



## Security Best Practices

- Enforce Pod Security Standards
- Default-deny Network Policies
- Image scanning (Trivy, Clair)
- External Secret Stores
- Admission Controllers
- API server audit logs